

# dog\_app

February 21, 2020

## 1 Convolutional Neural Networks

### 1.1 Project: Write an Algorithm for a Dog Identification App

---

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

**Note:** Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

## Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets: \* Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dogImages.

- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

*Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.*

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays `human_files` and `dog_files`.

```
In [9]: import numpy as np
        from glob import glob

        # load filenames for human and dog images
        human_files = np.array(glob("/data/lfw/**/*.jpg"))
        dog_files = np.array(glob("/data/dog_images/**/*.jpg"))

        # print number of images in each dataset
        print('There are %d total human images.' % len(human_files))
        print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

### ## Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the `haarcascades` directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [10]: import cv2
         import matplotlib.pyplot as plt
         %matplotlib inline

         # extract pre-trained face detector
         face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

         # load color (BGR) image
         img = cv2.imread(human_files[0])
         # convert BGR image to grayscale
         gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

         # find faces in image
         faces = face_cascade.detectMultiScale(gray)

         # print number of faces detected in the image
         print('Number of faces detected:', len(faces))

         # get bounding box for each detected face
         for (x,y,w,h) in faces:
             # add bounding box to color image
             cv2.rectangle(img, (x,y), (x+w,y+h), (255,0,0), 2)
```

```
# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()
```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

### 1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [11]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
```

```
img = cv2.imread(img_path)
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
faces = face_cascade.detectMultiScale(gray)
return len(faces) > 0
```

### 1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

**Question 1:** Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

**Answer:** (You can print out your results and/or write your percentages in this cell)

```
In [12]: from tqdm import tqdm
```

```
human_files_short = human_files[:100]
dog_files_short = dog_files[:100]
```

```
##-## Do NOT modify the code above this line. ##-##
```

```
## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
```

```
human_faces_recognized = [1 for i in human_files_short if face_detector(i)]
dog_faces_recognized = [1 for i in dog_files_short if face_detector(i)]
```

```
print('%{} of human faces detected.'.format(len(human_faces_recognized)/len(human_files_short)))
print('%{} of dog faces detected.'.format(len(dog_faces_recognized)/len(dog_files_short)))
```

```
%98.0 of human faces detected.
```

```
%17.0 of dog faces detected.
```

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [13]: ### (Optional)
```

```
### TODO: Test performance of another face detection algorithm.
```

```
### Feel free to use as many code cells as needed.
```

---

**## Step 2: Detect Dogs**

In this section, we use a [pre-trained model](#) to detect dogs in images.

### 1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of 1000 categories.

```
In [14]: import torch
import torchvision.models as models

# define VGG16 model
VGG16 = models.vgg16(pretrained=True)

# check if CUDA is available
use_cuda = torch.cuda.is_available()

# move model to GPU if CUDA is available
if use_cuda:
    VGG16 = VGG16.cuda()

#ME
print(VGG16)
print(VGG16.classifier[6].in_features)
print(VGG16.classifier[6].out_features)

VGG(
(features): Sequential(
  (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU(inplace)
  (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (3): ReLU(inplace)
  (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (6): ReLU(inplace)
  (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (8): ReLU(inplace)
  (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (11): ReLU(inplace)
  (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (13): ReLU(inplace)
  (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (15): ReLU(inplace)
  (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (18): ReLU(inplace)
  (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (20): ReLU(inplace)
```

```

(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace)
(23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(25): ReLU(inplace)
(26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(27): ReLU(inplace)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace)
(30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace)
  (2): Dropout(p=0.5)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace)
  (5): Dropout(p=0.5)
  (6): Linear(in_features=4096, out_features=1000, bias=True)
)
)
4096
1000

```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

### 1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher\_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```

In [28]: from PIL import Image
import torchvision.transforms as transforms

# Set PIL to be tolerant of image files that are truncated.
from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

def process_image(img_path):
    ''' Scales, crops, and normalizes a PIL image for a PyTorch model,
        returns an Numpy array
    '''

```

```

pil_im = Image.open(img_path, 'r')

pil_im.thumbnail((256,256))
pil_im=pil_im.crop((16,16,240,240))

#ish(np.asarray(pil_im))
np_image = np.array(pil_im)

#couldn't make it the other way
transformations = transforms.Compose([transforms.ToTensor(),transforms.Normalize((0
                                                                                   (0.229, 0.224, 0.225)))]

np_image = transformations(np_image).float()
return np_image

def VGG16_predict(img_path):
    '''
Use pre-trained VGG-16 model to obtain index corresponding to
predicted ImageNet class for image at specified path

Args:
    img_path: path to an image

Returns:
    Index corresponding to VGG-16 model's prediction
'''

    ## TODO: Complete the function.
    ## Load and pre-process an image from the given img_path
    ## Return the *index* of the predicted class for that image

    if use_cuda:
        VGG16.cuda()

    VGG16.eval()

    with torch.no_grad():

        image = process_image(img_path)
        if use_cuda:
            image = image.type(torch.FloatTensor).cuda()
            #I do not understand why we shold do this why not 3 dimensions are not sufficen
            #print(image.shape)
            image = image.unsqueeze(0)
            #print(image.shape)

        output = VGG16.forward(image)

```

```

        #ps = torch.exp(logps)
        _,cls_idx=torch.max(output, 1)
        #probs_tensor, classes_tensor = ps.topk(1,dim=1)

    #return probs_tensor, classes_tensor
    return cls_idx.item()

#imaj=process_image("a.jpg")
#imaj.shape
predict_that=VGG16_predict("/data/dog_images/test/022.Belgian_tervuren/Belgian_tervuren_01561.jpg")
print(predict_that)

```

224

### 1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the dog\_detector function below, which returns True if a dog is detected in an image (and False if not).

```

In [32]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    ## TODO: Complete the function.
    prediction=VGG16_predict(img_path)

    return not (prediction<151 or prediction>268)
dog_detector("/data/dog_images/test/022.Belgian_tervuren/Belgian_tervuren_01561.jpg")

```

Out[32]: True

### 1.1.6 (IMPLEMENTATION) Assess the Dog Detector

**Question 2:** Use the code cell below to test the performance of your dog\_detector function.

- What percentage of the images in human\_files\_short have a detected dog?
- What percentage of the images in dog\_files\_short have a detected dog?

**Answer:**

- %0.0 of human faces detected.
- %100.0 of dog faces detected.



```
In [9]: ### TODO: Test the performance of the dog_detector function
        ### on the images in human_files_short and dog_files_short.
        humano_faces_recognized = [1 for i in human_files_short if dog_detector(i)]
        dogo_faces_recognized = [1 for i in dog_files_short if dog_detector(i)]
        print('%{} of human faces detected.'.format(len(humano_faces_recognized)/len(human_files_short)))
        print('%{} of dog faces detected.'.format(len(dogo_faces_recognized)/len(dog_files_short)))

%0.0 of human faces detected.
%100.0 of dog faces detected.
```

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [10]: ### (Optional)
         ### TODO: Report the performance of another pre-trained network.
         ### Feel free to use as many code cells as needed.
```

---

### ## Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

### 1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dogImages/train`, `dogImages/valid`, and `dogImages/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [16]: import os
         from torchvision import datasets
         import torchvision.transforms as transforms
         ### TODO: Write data loaders for training, validation, and test sets
         ## Specify appropriate transforms, and batch_sizes
         data_dir = '/data/dog_images/'
         train_dir = os.path.join(data_dir, 'train/')
         valid_dir = os.path.join(data_dir, 'valid/')
         test_dir = os.path.join(data_dir, 'test/')
         train_transform = transforms.Compose([transforms.Resize(130), transforms.CenterCrop(128),
                                             transforms.RandomResizedCrop(128),
                                             transforms.RandomHorizontalFlip(),
                                             transforms.RandomRotation(50),
                                             transforms.ToTensor(),
                                             transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                                                    std=[0.229, 0.224, 0.225])])

         data_transform = transforms.Compose([transforms.Resize(130), transforms.CenterCrop(128),
                                             transforms.ToTensor(),
                                             transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                                                    std=[0.229, 0.224, 0.225])])

         train_data = datasets.ImageFolder(train_dir, transform=train_transform)
         valid_data = datasets.ImageFolder(valid_dir, transform=data_transform)
         test_data = datasets.ImageFolder(test_dir, transform=data_transform)

         #----I've Just added this to understand what is going on-----
```

```

batch_size = 20
num_workers=0

# prepare data loaders
train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                             num_workers=num_workers, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
                                           num_workers=num_workers, shuffle=True)
valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=batch_size,
                                           num_workers=num_workers, shuffle=True)

loaders_scratch={}
loaders_scratch["train"]=train_loader
loaders_scratch["valid"]=valid_loader
loaders_scratch["test"]=test_loader

dataiter = iter(train_loader)
images, labels = dataiter.next()

#unnormalizing images for display
def im_convert(tensor):
    """ Display a tensor as an image. """

    image = tensor.to("cpu").clone().detach()
    image = image.numpy().squeeze()

    image = image.transpose(1,2,0)
    image = image * np.array((0.229, 0.224, 0.225)) + np.array((0.485, 0.456, 0.406))
    image = image.clip(0, 1)

    return image

fig = plt.figure(figsize=(25, 4))
# plot the images in the batch, along with the corresponding labels
for idx in np.arange(20):
    c_image=im_convert(images[idx])
    ax = fig.add_subplot(2, 20/2, idx+1, xticks=[], yticks=[])
    plt.imshow(c_image)
#----I've Just added this to understand what is going on-----

```



**Question 3:** Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

**Answer:** - I resized images on the dataset via torchvision.transform tool bu cropping. there were various sizes of images in the dataset that some of them are very big so 224 seemed reasonable at this time. - Yes, again with the help of transform tool.I made them(only training portion) flipped,cropped,rotated.

### 1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```
In [17]: import torch.nn as nn
import torch.nn.functional as F

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN

        self.conv1 = nn.Conv2d(3, 16, 3, padding=1)

        self.conv2 = nn.Conv2d(16, 32, 3, padding=1)

        self.conv3 = nn.Conv2d(32, 64, 3, padding=1)
        # max pooling layer
        self.pool = nn.MaxPool2d(2, 2)

        #my image size is 128*128, so after third pool layer it becomes 16*16
        self.fc1 = nn.Linear(64 * 16 * 16, 500)
        # linear layer (500 -> 133)
        self.fc2 = nn.Linear(500, 133)
        # dropout layer (p=0.25)
        self.dropout = nn.Dropout(0.25)

    def forward(self, x):
        ## Define forward behavior
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = self.pool(F.relu(self.conv3(x)))
        # flatten image input
        x = x.view(-1, 64 * 16 * 16)
```

```

        # add dropout layer
        x = self.dropout(x)
        # add 1st hidden layer, with relu activation function
        x = F.relu(self.fc1(x))
        # add dropout layer
        x = self.dropout(x)
        # add 2nd hidden layer, with relu activation function
        x = self.fc2(x)
        return x

### You do NOT have to modify the code below this line. ###

# instantiate the CNN
model_scratch = Net()

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

```

**Question 4:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

**Answer:** - I start with 128\*128 image as an input to keep model size a bit lower compare to VGG16, - create three Convolutional layer sequentially as feature extractors that double the feature maps each step, - then applied relu function for adding nonlinearity - finally applied pooling operation to lower the number of paramaters and discard the spatial information. - last two layers for classification. Just before first one, I flatten out all the information come from feature maps. - For last layer I have 133 output for each dog breeds - To prevent overfitting add droppout layer for FC layers

### 1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```

In [19]: import torch.optim as optim

        ### TODO: select loss function
        criterion_scratch = nn.CrossEntropyLoss()

        ### TODO: select optimizer
        optimizer_scratch = optim.SGD(model_scratch.parameters(), lr=0.01)

```

### 1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```

In [20]: # the following import is required for training to be robust to truncated images
        from PIL import ImageFile

```

```

ImageFile.LOAD_TRUNCATED_IMAGES = True

def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        #####
        # train the model #
        #####
        model.train()
        for batch_idx, (data, target) in enumerate(loaders['train']):
            #for batch_idx, (data, target) in enumerate(train_loader):
                # move to GPU
                if use_cuda:
                    data, target = data.cuda(), target.cuda()

            optimizer.zero_grad()

            output = model(data)

            loss = criterion(output, target)

            loss.backward()
            optimizer.step()
            ## find the loss and update the model parameters accordingly
            ## record the average training loss, using something like
            ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
            train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
        #####
        # validate the model #
        #####
        model.eval()
        for batch_idx, (data, target) in enumerate(loaders['valid']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            ## update the average validation loss
            # forward pass: compute predicted outputs by passing inputs to the model
            output = model(data)
            # calculate the batch loss
            loss = criterion(output, target)

```

```

        # update average validation loss
        valid_loss += loss.item()*data.size(0)
    # print training/validation statistics
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
        epoch,
        train_loss,
        valid_loss
    ))

    ## TODO: save the model if validation loss has decreased
    if valid_loss <= valid_loss_min:
        print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.fo
            valid_loss_min,
            valid_loss))
        torch.save(model.state_dict(), save_path)
        valid_loss_min = valid_loss

    # return trained model
    return model

# train the model
model_scratch = train(40, loaders_scratch, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, 'model_scratch.pt')

#load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))

```

```

Epoch: 1      Training Loss: 4.883419      Validation Loss: 4066.713603
Validation loss decreased (inf --> 4066.713603). Saving model ...
Epoch: 2      Training Loss: 4.858892      Validation Loss: 4040.409513
Validation loss decreased (4066.713603 --> 4040.409513). Saving model ...
Epoch: 3      Training Loss: 4.792527      Validation Loss: 3951.604486
Validation loss decreased (4040.409513 --> 3951.604486). Saving model ...
Epoch: 4      Training Loss: 4.664647      Validation Loss: 3824.049020
Validation loss decreased (3951.604486 --> 3824.049020). Saving model ...
Epoch: 5      Training Loss: 4.579844      Validation Loss: 3768.377526
Validation loss decreased (3824.049020 --> 3768.377526). Saving model ...
Epoch: 6      Training Loss: 4.538202      Validation Loss: 3743.325334
Validation loss decreased (3768.377526 --> 3743.325334). Saving model ...
Epoch: 7      Training Loss: 4.515414      Validation Loss: 3720.292847
Validation loss decreased (3743.325334 --> 3720.292847). Saving model ...
Epoch: 8      Training Loss: 4.495499      Validation Loss: 3686.268785
Validation loss decreased (3720.292847 --> 3686.268785). Saving model ...
Epoch: 9      Training Loss: 4.473136      Validation Loss: 3694.879448
Epoch: 10     Training Loss: 4.450355      Validation Loss: 3669.316256
Validation loss decreased (3686.268785 --> 3669.316256). Saving model ...
Epoch: 11     Training Loss: 4.432211      Validation Loss: 3711.634200
Epoch: 12     Training Loss: 4.404918      Validation Loss: 3625.468256

```

Validation loss decreased (3669.316256 --> 3625.468256). Saving model ...  
 Epoch: 13            Training Loss: 4.380626            Validation Loss: 3631.659245  
 Epoch: 14            Training Loss: 4.364997            Validation Loss: 3588.462188  
 Validation loss decreased (3625.468256 --> 3588.462188). Saving model ...  
 Epoch: 15            Training Loss: 4.341889            Validation Loss: 3588.553038  
 Epoch: 16            Training Loss: 4.334906            Validation Loss: 3573.654652  
 Validation loss decreased (3588.462188 --> 3573.654652). Saving model ...  
 Epoch: 17            Training Loss: 4.301415            Validation Loss: 3582.538617  
 Epoch: 18            Training Loss: 4.271464            Validation Loss: 3536.978762  
 Validation loss decreased (3573.654652 --> 3536.978762). Saving model ...  
 Epoch: 19            Training Loss: 4.274193            Validation Loss: 3503.114505  
 Validation loss decreased (3536.978762 --> 3503.114505). Saving model ...  
 Epoch: 20            Training Loss: 4.240935            Validation Loss: 3666.617088  
 Epoch: 21            Training Loss: 4.234035            Validation Loss: 3542.543993  
 Epoch: 22            Training Loss: 4.200685            Validation Loss: 3597.661912  
 Epoch: 23            Training Loss: 4.207452            Validation Loss: 3439.505157  
 Validation loss decreased (3503.114505 --> 3439.505157). Saving model ...  
 Epoch: 24            Training Loss: 4.184228            Validation Loss: 3460.990825  
 Epoch: 25            Training Loss: 4.151198            Validation Loss: 3449.332516  
 Epoch: 26            Training Loss: 4.126995            Validation Loss: 3432.598786  
 Validation loss decreased (3439.505157 --> 3432.598786). Saving model ...  
 Epoch: 27            Training Loss: 4.103266            Validation Loss: 3453.349491  
 Epoch: 28            Training Loss: 4.088383            Validation Loss: 3406.095228  
 Validation loss decreased (3432.598786 --> 3406.095228). Saving model ...  
 Epoch: 29            Training Loss: 4.077822            Validation Loss: 3380.612673  
 Validation loss decreased (3406.095228 --> 3380.612673). Saving model ...  
 Epoch: 30            Training Loss: 4.043926            Validation Loss: 3453.760935  
 Epoch: 31            Training Loss: 4.049820            Validation Loss: 3361.085839  
 Validation loss decreased (3380.612673 --> 3361.085839). Saving model ...  
 Epoch: 32            Training Loss: 4.017241            Validation Loss: 3361.049440  
 Validation loss decreased (3361.085839 --> 3361.049440). Saving model ...  
 Epoch: 33            Training Loss: 3.985492            Validation Loss: 3322.024257  
 Validation loss decreased (3361.049440 --> 3322.024257). Saving model ...  
 Epoch: 34            Training Loss: 3.969734            Validation Loss: 3447.742255  
 Epoch: 35            Training Loss: 3.942060            Validation Loss: 3301.541464  
 Validation loss decreased (3322.024257 --> 3301.541464). Saving model ...  
 Epoch: 36            Training Loss: 3.960695            Validation Loss: 3296.066978  
 Validation loss decreased (3301.541464 --> 3296.066978). Saving model ...  
 Epoch: 37            Training Loss: 3.918153            Validation Loss: 3333.022668  
 Epoch: 38            Training Loss: 3.902590            Validation Loss: 3269.852604  
 Validation loss decreased (3296.066978 --> 3269.852604). Saving model ...  
 Epoch: 39            Training Loss: 3.891379            Validation Loss: 3322.333426  
 Epoch: 40            Training Loss: 3.887197            Validation Loss: 3228.919644  
 Validation loss decreased (3269.852604 --> 3228.919644). Saving model ...



### 1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
In [21]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

    # call test function
    test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)
```

Test Loss: 3.900254

Test Accuracy: 12% (104/836)

---

## Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)  
You will now use transfer learning to create a CNN that can identify dog breed from images.  
Your CNN must attain at least 60% accuracy on the test set.

### 1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```
In [22]: ## TODO: Specify data loaders
data_dir = '/data/dog_images/'
train_dir = os.path.join(data_dir, 'train/')
valid_dir = os.path.join(data_dir, 'valid/')
test_dir = os.path.join(data_dir, 'test/')
train_transform = transforms.Compose([transforms.Resize(250), transforms.CenterCrop(224),
                                     transforms.RandomResizedCrop(224),
                                     transforms.RandomHorizontalFlip(),
                                     transforms.RandomRotation(50),
                                     transforms.ToTensor()])
""" , transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)) """

data_transform = transforms.Compose([transforms.Resize(250), transforms.CenterCrop(224),
                                     transforms.ToTensor()])
""" , transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)) """

train_data = datasets.ImageFolder(train_dir, transform=train_transform)
valid_data = datasets.ImageFolder(valid_dir, transform=data_transform)
test_data = datasets.ImageFolder(test_dir, transform=data_transform)

#----I've Just added this to understand what is going on-----

batch_size = 20
num_workers=0

# prepare data loaders
train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                           num_workers=num_workers, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
                                           num_workers=num_workers, shuffle=True)
valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=batch_size,
                                           num_workers=num_workers, shuffle=True)

loaders_transfer={}
loaders_transfer["train"]=train_loader
loaders_transfer["valid"]=valid_loader
loaders_transfer["test"]=test_loader
```

### 1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```
In [23]: import torchvision.models as models
import torch.nn as nn

## TODO: Specify model architecture
model_transfer = models.vgg16(pretrained=True)
print(model_transfer)
for param in model_transfer.features.parameters():
    param.requires_grad = False
#number of dog classes in the dataset
dog_classes=133
n_inputs = model_transfer.classifier[6].in_features

# new layers automatically have requires_grad = True
last_layer = nn.Linear(n_inputs, dog_classes)
model_transfer.classifier[6] = last_layer
print(model_transfer)
# check if CUDA is available
use_cuda = torch.cuda.is_available()

# move model to GPU if CUDA is available
if use_cuda:
    model_transfer = model_transfer.cuda()

VGG(
(features): Sequential(
  (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU(inplace)
  (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (3): ReLU(inplace)
  (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (6): ReLU(inplace)
  (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (8): ReLU(inplace)
  (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (11): ReLU(inplace)
  (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (13): ReLU(inplace)
  (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (15): ReLU(inplace)
```

```

(16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(18): ReLU(inplace)
(19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace)
(23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(25): ReLU(inplace)
(26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(27): ReLU(inplace)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace)
(30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace)
  (2): Dropout(p=0.5)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace)
  (5): Dropout(p=0.5)
  (6): Linear(in_features=4096, out_features=1000, bias=True)
)
)
VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace)
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (13): ReLU(inplace)
    (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (15): ReLU(inplace)
    (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (18): ReLU(inplace)
    (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

(20): ReLU(inplace)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace)
(23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(25): ReLU(inplace)
(26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(27): ReLU(inplace)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace)
(30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace)
  (2): Dropout(p=0.5)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace)
  (5): Dropout(p=0.5)
  (6): Linear(in_features=4096, out_features=133, bias=True)
)
)

```

**Question 5:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

**Answer:** - I start with a pre-trained model VGG16 as a feature extractor. - to use its pre trained weights freeze them before training. - I replaced current last fully connected layer with new one that is more suitable with my 133 dog classes. - This VGG16 model usable for my needs because it trained on a huge dataset that also includes most of my dog breeds. So features it find would be compatible with my dataset. Only additional step is training fully connected layers.

#### 1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```

In [24]: import torch.optim as optim
         criterion_transfer = nn.CrossEntropyLoss()
         optimizer_transfer = optim.SGD(model_transfer.classifier.parameters(), lr=0.01)

```

#### 1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```

In [25]: # train the model
         n_epochs=20
         model_transfer = train(n_epochs, loaders_transfer, model_transfer, optimizer_transfer,

```

```
# load the model that got the best validation accuracy (uncomment the line below)
#model_transfer.load_state_dict(torch.load('model_transfer.pt'))
```

```
Epoch: 1      Training Loss: 4.206079      Validation Loss: 1788.600531
Validation loss decreased (inf --> 1788.600531). Saving model ...
Epoch: 2      Training Loss: 3.237751      Validation Loss: 1261.074567
Validation loss decreased (1788.600531 --> 1261.074567). Saving model ...
Epoch: 3      Training Loss: 2.935025      Validation Loss: 1150.476137
Validation loss decreased (1261.074567 --> 1150.476137). Saving model ...
Epoch: 4      Training Loss: 2.747016      Validation Loss: 1010.027809
Validation loss decreased (1150.476137 --> 1010.027809). Saving model ...
Epoch: 5      Training Loss: 2.651780      Validation Loss: 946.398115
Validation loss decreased (1010.027809 --> 946.398115). Saving model ...
Epoch: 6      Training Loss: 2.591071      Validation Loss: 975.882382
Epoch: 7      Training Loss: 2.508532      Validation Loss: 924.886099
Validation loss decreased (946.398115 --> 924.886099). Saving model ...
Epoch: 8      Training Loss: 2.458711      Validation Loss: 909.383541
Validation loss decreased (924.886099 --> 909.383541). Saving model ...
Epoch: 9      Training Loss: 2.426657      Validation Loss: 903.601698
Validation loss decreased (909.383541 --> 903.601698). Saving model ...
Epoch: 10     Training Loss: 2.393345      Validation Loss: 873.990644
Validation loss decreased (903.601698 --> 873.990644). Saving model ...
Epoch: 11     Training Loss: 2.326861      Validation Loss: 852.366130
Validation loss decreased (873.990644 --> 852.366130). Saving model ...
Epoch: 12     Training Loss: 2.313775      Validation Loss: 841.157330
Validation loss decreased (852.366130 --> 841.157330). Saving model ...
Epoch: 13     Training Loss: 2.281017      Validation Loss: 851.368233
Epoch: 14     Training Loss: 2.203907      Validation Loss: 871.683766
Epoch: 15     Training Loss: 2.238973      Validation Loss: 851.991068
Epoch: 16     Training Loss: 2.200655      Validation Loss: 822.150648
Validation loss decreased (841.157330 --> 822.150648). Saving model ...
Epoch: 17     Training Loss: 2.219464      Validation Loss: 806.875275
Validation loss decreased (822.150648 --> 806.875275). Saving model ...
Epoch: 18     Training Loss: 2.158723      Validation Loss: 779.548876
Validation loss decreased (806.875275 --> 779.548876). Saving model ...
Epoch: 19     Training Loss: 2.165520      Validation Loss: 809.705009
Epoch: 20     Training Loss: 2.144737      Validation Loss: 856.258096
```

### 1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [26]: model_transfer.load_state_dict(torch.load('model_transfer.pt'))
         test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

Test Loss: 0.938510

Test Accuracy: 71% (600/836)

### 1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [29]: ### TODO: Write a function that takes a path to an image as input
         ### and returns the dog breed that is predicted by the model.
         #loaders_transfer["train"]

# list of class names by index, i.e. a name can be accessed like class_names[0]
#this is original
#class_names = [item[4:].replace("_", " ") for item in data_transfer['train'].classes]
class_names = [item[4:].replace("_", " ") for item in train_data.classes]

def predict_breed_transfer(img_path):

    if use_cuda:
        model_transfer.cuda()

    model_transfer.eval()

    with torch.no_grad():

        image = process_image(img_path)
        if use_cuda:
            image = image.type(torch.FloatTensor).cuda()

        image = image.unsqueeze(0)
        #print(image.shape)

        output = model_transfer.forward(image)

        #ps = torch.exp(logps)
        _,cls_idx=torch.max(output, 1)
        #probs_tensor, classes_tensor = ps.topk(1,dim=1)
```



Sample Human Output

```
#return probs_tensor, classes_tensor
return class_names[cls_idx.item()]
predict_that=predict_breed_transfer("/data/dog_images/test/022.Belgian_tervuren/Belgian
#predict_that=predict_breed_transfer("/data/dog_images/test/001.Affenpinscher/Affenpins
print(predict_that)
```

Belgian terfvuren

#### ## Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `dog_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

#### 1.1.18 (IMPLEMENTATION) Write your Algorithm

In [30]: *### TODO: (Write your algorithm.*  
*### Feel free to use as many code cells as needed.*

```
def run_app(img_path):
    print("\n\n")
    print("file name: {}".format(img_path))
    print("-----")
    if face_detector(img_path):
        print("This human is look like:")
```



```

img = cv2.imread(img_path)
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
# find faces in image
faces = face_cascade.detectMultiScale(gray)
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
plt.imshow(cv_rgb)
plt.show()
dog_like=predict_breed_transfer(img_path)
return print(dog_like)
if dog_detector(img_path):
    d_breed=predict_breed_transfer(img_path)
    print("This good dog's breed:")
    im=process_image(img_path)
    im=im_convert(im)
    plt.imshow(im)
    plt.show()
    return print(d_breed)
else:
    print("no humans or dogs at this time")

## handle cases for a human face, dog, and neither
aa=run_app("/data/lfw/Alex_Corretja/Alex_Corretja_0001.jpg")

```

---

### ## Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

#### 1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

**Question 6:** Is the output better than you expected :) ? Or worse :( ? Provide at least three possible points of improvement for your algorithm.

**Answer:** - My algorithm doesn't check for more than one person or different dog breeds in the image. - Normalizing the input image could be helpful. - Tuning VGG16 for not 1000 categories but only dog categories could be helpful. - There could be a option for use of Cuda or not. - My images are slowing the app. because they aren't resized.

```

In [33]: ## TODO: Execute your algorithm from Step 6 on
## at least 6 images on your computer.
## Feel free to use as many code cells as needed.
my_human_files = np.array(glob("my_human/*"))
my_dog_files = np.array(glob("my_dog/*"))

```

```


## suggested code, below
for file in np.hstack((human_files[:3], dog_files[:3])):
    run_app(file)

for my_file in np.hstack((my_human_files, my_dog_files)):
    run_app(my_file)
print('There are %d total human images.' % len(my_human_files))
print('There are %d total dog images.' % len(my_dog_files))

```

file name: /data/lfw/Dan\_Ackroyd/Dan\_Ackroyd\_0001.jpg


-----  
This human is look like:



Ibizan hound

file name: /data/lfw/Alex\_Corretja/Alex\_Corretja\_0001.jpg


-----  
This human is look like:



Bearded collie

file name: /data/lfw/Daniele\_Bergamin/Daniele\_Bergamin\_0001.jpg

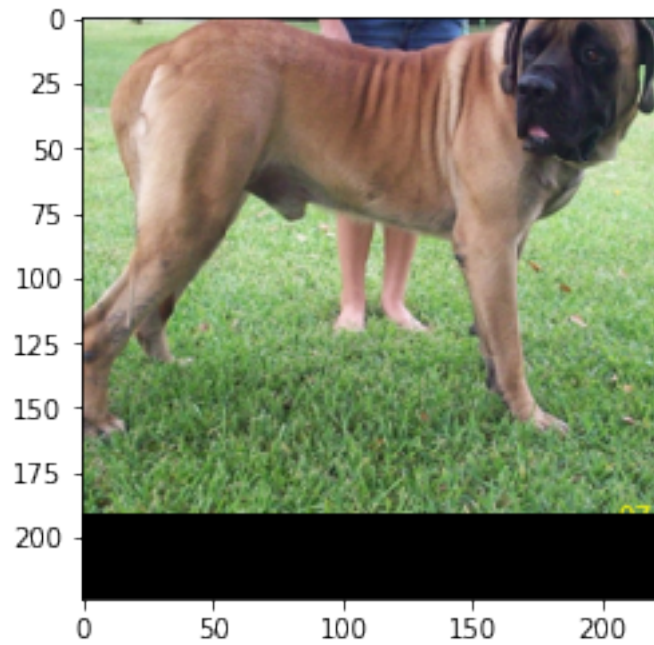
-----  
This human is look like:



German pinscher

file name: /data/dog\_images/train/103.Mastiff/Mastiff\_06833.jpg


-----  
This good dog's breed:



Bullmastiff

file name: /data/dog\_images/train/103.Mastiff/Mastiff\_06826.jpg

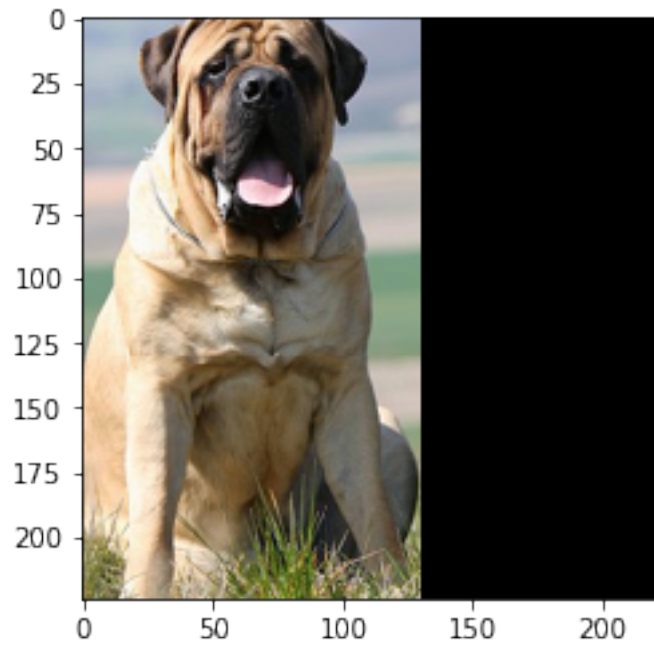
-----  
This good dog's breed:



Bullmastiff

file name: /data/dog\_images/train/103.Mastiff/Mastiff\_06871.jpg


-----  
This good dog's breed:



Bullmastiff

file name: my\_human/2016-01-05 13.00.48.jpg

-----  
This human is look like:




English toy spaniel

file name: my\_human/2015-11-13 13.41.54.jpg

-----  
This human is look like:






Dalmatian

file name: my\_human/2015-12-31 14.08.45.jpg

-----  
no humans or dogs at this time

file name: my\_human/2015-11-08 14.31.54.jpg


-----  
This human is look like:



Afghan hound

file name: my\_human/2016-01-12 03.10.32.jpg


-----  
This human is look like:



Poodle

file name: my\_human/\_MG\_9755.JPG

-----  
This human is look like:



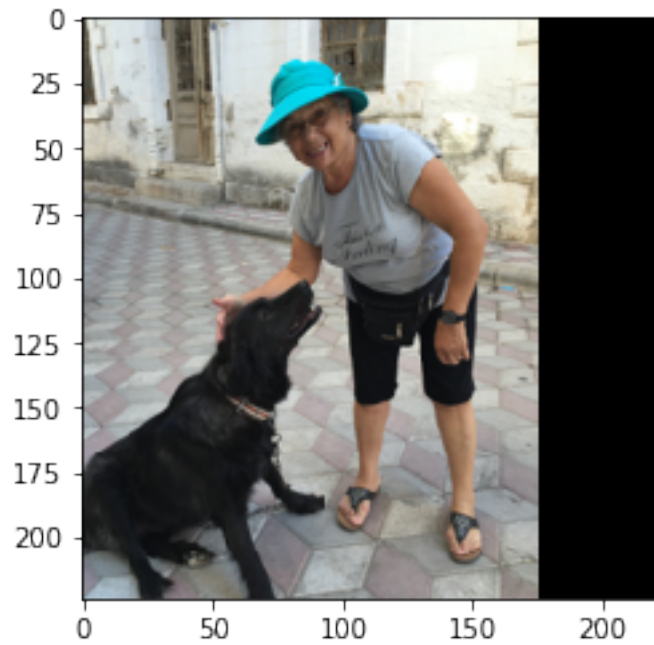
Dogue de bordeaux

file name: my\_dog/2016-02-02 11.47.53.jpg

-----  
no humans or dogs at this time

file name: my\_dog/2015-09-01 16.58.55.jpg

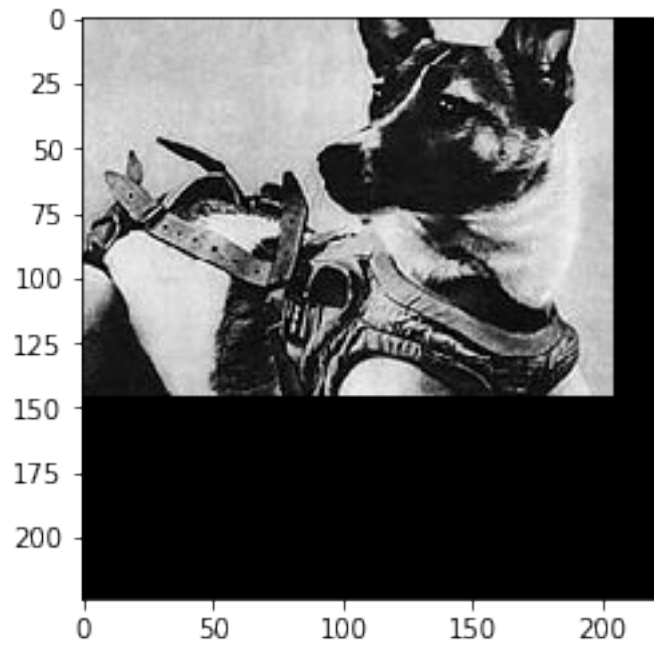
-----  
This good dog's breed:



Giant schnauzer

file name: my\_dog/2015-12-14 12.39.01.jpg

-----  
This good dog's breed:




Norwegian elkhound

file name: my\_dog/2016-02-02 11.50.09.jpg

-----  
no humans or dogs at this time

file name: my\_dog/20190526\_163123269\_iOS.jpg

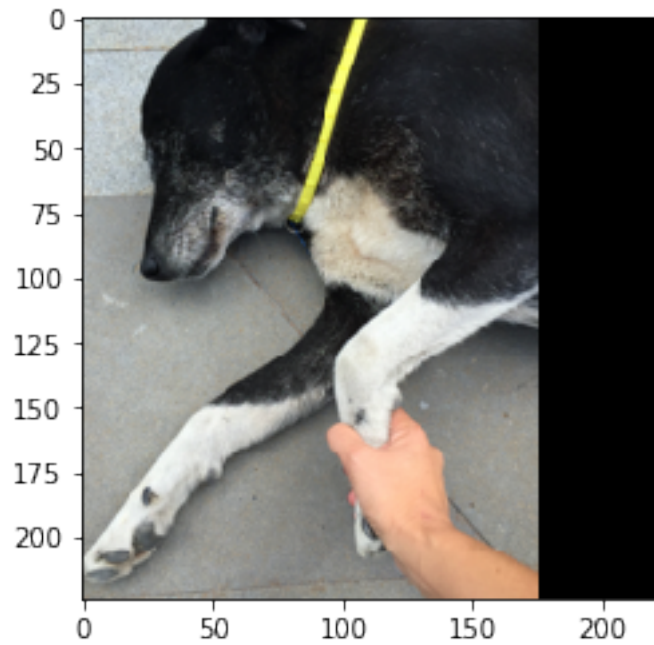
-----  
This human is look like:



Parson russell terrier

file name: my\_dog/2015-09-15 15.26.25.jpg

-----  
This good dog's breed:



Australian cattle dog

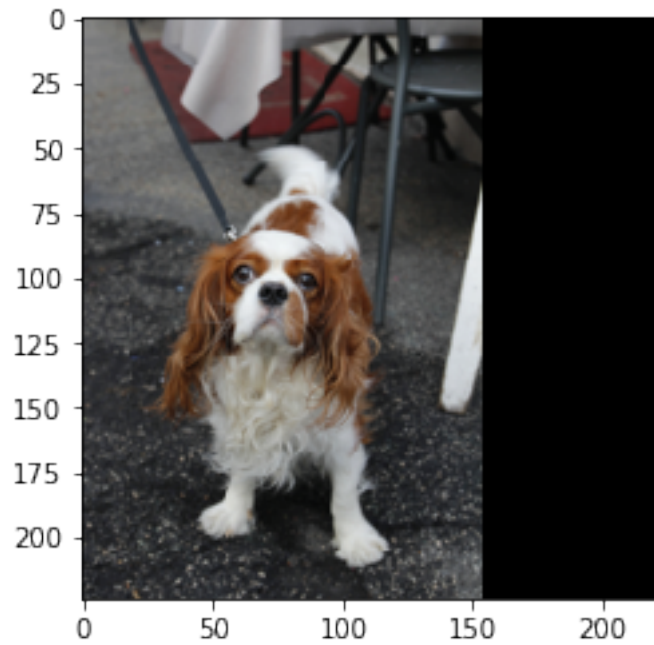
file name: my\_dog/2016-02-02 11.51.15.jpg

-----  
no humans or dogs at this time

file name: my\_dog/2016-02-01 10.45.02.jpg

-----  
This good dog's breed:






Cavalier king charles spaniel

file name: my\_dog/20190526\_163620251\_i0S.jpg

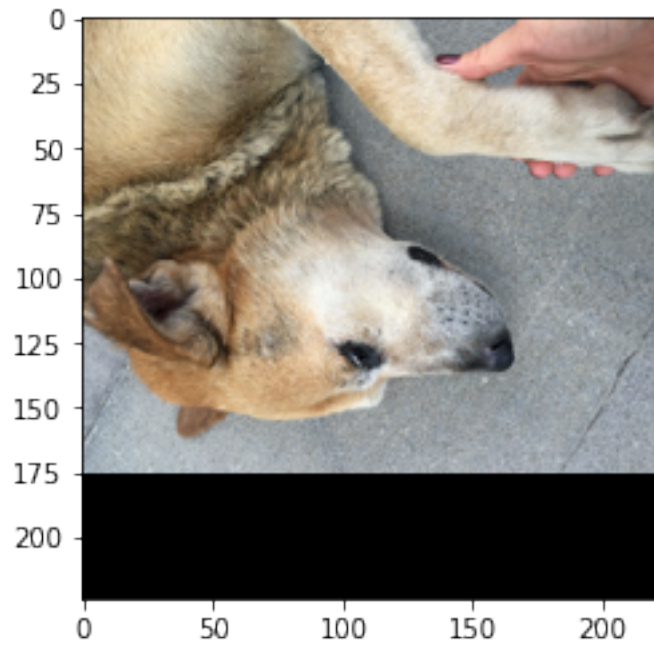
-----  
This human is look like:



Belgian malinois

file name: my\_dog/2015-09-15 15.26.00.jpg

-----  
This good dog's breed:



Border terrier

file name: my\_dog/20190508\_164639998\_iOS.jpg

-----  
no humans or dogs at this time  
There are 6 total human images.  
There are 11 total dog images.

In [ ]: